

# CS61C Summer 2018 Project 3: CPU

TA: Sruthi Veeragandham

Due July 20, 2018 @ 11:59:59 pm

NOTE: All important updates will be posted on the Piazza Thread titled "Project 3 RELEASED, Updates Here". Please prioritize Piazza as the source of truth and check it regularly.

## IMPORTANT INFO – PLEASE READ

- Make sure that you check the Piazza announcement about the skeleton code bug in `cpu.circ` if you started working before 9:30PM on Saturday 7/14. If not, you don't have to worry about the bug, as the starter has been updated.
- **USE STEPHAN'S LOGISIM-EVOLUTION FIX FROM PIAZZA, NOT THE LOGISIM-EVOLUTION THAT CAME WITH LAB 5 OR THE LOGISIM-EVOLUTION FROM THE LOGISIM-EVOLUTION GITHUB.** [The correct logisim is linked here.](#) The correct logisim-evolution is included with the proj3 starter code, too. This should fix the issue many of you had with getting an error about X11 forwarding on your local machine. (The only difference between this .jar file and the one you have been using is the fix to the X11 issue, so all your tests should still pass if they were passing and all your circuits should look exactly the same).
- You are allowed to use any of Logisim's built-in blocks for all parts of this project.
- **Save often.** Logisim can be buggy and the last thing you want is to lose some of your hard work. There are students every semester who have had to start over large chunks of their projects due to this.
- Approach this project like you would any coding assignment: construct it piece by piece and test each component early and often!
- **Make sure to check that `cpu.circ` fits into `run.circ`.** (This also means that you should not be moving around given **inputs** and **outputs** ports in the circuit files `alu.circ`, `regfile.circ`, `mem.circ`, and `cpu.circ`).
- Because the files you are working on are not plain code and circuit schematics, they can't really be merged. **Do not work on the file in two places and attempt to merge! It will not work, you will be sad, and I will be sad too.**
- **Make sure all your changes are located in `cpu.circ`.** If you need to modularize, use sub-circuits instead of a new library module.

## Overview

In this project you will be using [logisim-evolution](#) to implement a 32-bit two-cycle processor based on RISC-V. This project is meant to give you a better understanding of the actual RISC-V datapath. In fact, after this project you will know everything you need in order to build a RISC-V CPU in Logisim that could understand your assembled and linked input from Project 2!

You will be completing a 2-stage pipelined processor.

## 0) Obtaining the Files

We will provide you with skeleton code that includes sanity tests along with a CPU template (`cpu.circ`), harness (`run.circ`), and the data memory module (`mem.circ`). You will obtain the skeleton code from Github. You should also make a private BitBucket repo called proj3-xxx, where xxx is your 3-letter login. Make sure you give cs61c-staff admin access to your repo.

```
mkdir proj3
cd proj3
git init
git remote add proj3-starter https://github.com/61c-teach/su18-proj3-starter
git fetch proj3-starter
git merge proj3-starter/master -m "getting proj3 skeleton code"
git remote add origin https://bitbucket.org/your-username/proj3-xxx.git
git add -A
git commit -m "initial commit to personal repo"
git push origin master
```

Before you start, you need to complete the ALU and RegFile in Lab 6. After you have completed lab 6, copy your *alu.circ* and *regfile.circ* into your proj3 directory. Currently, the proj3 directory contains the empty *alu.circ* and *regfile.circ* skeleton code.

## 1) Getting Started – Processor

We have provided a skeleton for your processor in *cpu.circ*. You will be using your own implementations of the ALU and RegFile as you construct your datapath. If you find any errors in either of these components, you should definitely fix them. You are responsible for constructing the entire datapath and control from scratch. Your completed processor should implement the ISA detailed below in the section Instruction Set Architecture (ISA) using a two-cycle pipeline, specified below.

Your processor will get its program from the processor harness *run.circ*. Your processor will output the address of an instruction, and accept the instruction at that address as an input. Inspect *run.circ* to see exactly what's going on. (This same harness will be used to test your final submission, so make sure your CPU fits in the harness before submitting your work!) Your processor has 2 inputs that come from the harness:

INPUT NAME	BIT WIDTH	DESCRIPTION
INSTRUCTION	32	Driven with the instruction at the instruction memory address identified by the <code>FETCH_ADDRESS</code> (see below).
CLOCK	1	The input for the clock. As with the register file, this can be sent into subcircuits (e.g. the CLK input for your register file) or attached directly to the clock inputs of memory units in Logisim, but should not otherwise be gated (i.e., do not invert it, do not AND it with anything, etc.).

Your processor must provide the following outputs to the harness:

OUTPUT NAME	BIT WIDTH	DESCRIPTION
ra	32	Driven with the contents of ra. FOR TESTING
sp	32	Driven with the contents of sp. FOR TESTING
t0	32	Driven with the contents of t0. FOR TESTING
t1	32	Driven with the contents of t1. FOR TESTING
t2	32	Driven with the contents of t2. FOR TESTING
s0	32	Driven with the contents of s0. FOR TESTING
s1	32	Driven with the contents of s1. FOR TESTING
a0	32	Driven with the contents of a0. FOR TESTING
FETCH_ADDRESS	32	This output is used to select which instruction is presented to the processor on the INSTRUCTION input.

Just like in Lab 6, be careful **NOT** to move the input or output pins! You should ensure that your processor is correctly loaded by a fresh copy of *run.circ* before you submit. You can download a fresh copy from the starter repo website.

It is highly recommended that you start out by implementing your PC logic, so that you are generating the correct `FETCH_ADDR` output. Then, you can work through the instructions systematically, writing tests that cover the instructions you have implemented so far to confirm that you are on the right path (see the "Testing" section for more details on how to test). For a walkthrough of how to implement your first instruction, see the [guide in Exercise 3 of Lab 6](#).

### 1.5) Getting Started – Memory

The memory unit is already fully implemented for you! Here's a quick summary of its inputs and outputs:

OUTPUT NAME	IN- OR OUT-PUT?	BIT WIDTH	DESCRIPTION
A: ADDR	In	32	Address to read/write to in Memory
D: WRITE DATA	In	32	Value to be written to Memory
En: WRITE ENABLE	In	1	Equal to one on any instructions that write to memory, and zero otherwise
Clock	In	1	Driven by the clock input to <i>cpu.circ</i>
D: READ DATA	Out	32	Driven by the data stored at the specified address.

Note that the memory is word-addressable, meaning given an address, it will return 4-byte data.

## 2) The Instruction Set Architecture

Your CPU will support the instructions listed below. Most of the instructions should behave the same as the RISC-V you learned in class. Notice that this is not all of the instructions on the green card! These are the only instructions you will be tested on implementing, but you will not be punished for implementing other instructions (the autograder will cover ONLY the instructions below). If anything surprises you, it is likely that I made a mistake. Please make a Piazza post about it. As in Lab 6, div and rem should be unsigned. For mulh, the output we expected is the upper 32 bits of your result from the built-in Logisim multiply block, regardless of whether this matches with what you see in Venus (this will not necessarily match what you see in Venus if you are multiplying a negative number by a positive number; for most numbers, not all of the overflow bits will be needed and as a result the upper bits will not reflect the correct sign). If you pass the ALU tests in Lab 6, you should feel confident your ALU is working correctly. **Extra credit opportunity: if you implement mulh so that the output *\*does\** match what you see in Venus in all cases, you will receive some extra credit. This is not something you should work on until you have finished everything else.**

### Instruction Formats

31	30	25 24	21	20	19	15 14	12 11	8	7	6	0	
funct7		rs2			rs1	funct3	rd		opcode		R-type	
imm[11:0]					rs1	funct3	rd		opcode		I-type	
imm[11:5]		rs2	rs1	funct3	imm[4:0]		opcode		S-type			
imm[12]	imm[10:5]		rs2	rs1	funct3	imm[4:1]	imm[11]	opcode		B-type		
imm[31:12]							rd		opcode		U-type	
imm[20]	imm[10:1]		imm[11]	imm[19:12]		rd		opcode		J-type		

### The Instructions

INSTRUCTION	TYPE	OPCODE	FUNCT3	FUNCT7/IMM	OPERATION
add rd, rs1, rs2	R	0x33	0x0	0x00	$R[rd] \leftarrow R[rs1] + R[rs2]$
mul rd, rs1, rs2			0x0	0x01	$R[rd] \leftarrow (R[rs1] * R[rs2])[31:0]$
sub rd, rs1, rs2			0x0	0x20	$R[rd] \leftarrow R[rs1] - R[rs2]$
sll rd, rs1, rs2			0x1	0x00	$R[rd] \leftarrow R[rs1] \ll R[rs2]$
mulh rd, rs1, rs2			0x1	0x01	$R[rd] \leftarrow (R[rs1] * R[rs2])[63:32]$
slt rd, rs1, rs2			0x2	0x00	$R[rd] \leftarrow (R[rs1] < R[rs2]) ? 1 : 0$ (signed)
xor rd, rs1, rs2			0x4	0x00	$R[rd] \leftarrow R[rs1] \wedge R[rs2]$
div rd, rs1, rs2			0x4	0x01	(unsigned) $R[rd] \leftarrow R[rs1] / R[rs2]$
srl rd, rs1, rs2			0x5	0x00	$R[rd] \leftarrow R[rs1] \gg R[rs2]$
or rd, rs1, rs2			0x6	0x00	$R[rd] \leftarrow R[rs1] \vee R[rs2]$
rem rd, rs1, rs2			0x6	0x01	(unsigned) $R[rd] \leftarrow R[rs1] \% R[rs2]$
and rd, rs1, rs2			0x7	0x00	$R[rd] \leftarrow R[rs1] \& R[rs2]$
lb rd, offset(rs1)	I	0x03	0x0		$R[rd] \leftarrow \text{SignExt}(\text{Mem}(R[rs1] + \text{offset}, \text{byte}))$
lh rd, offset(rs1)			0x1		$R[rd] \leftarrow \text{SignExt}(\text{Mem}(R[rs1] + \text{offset}, \text{half}))$
lw rd, offset(rs1)			0x2		$R[rd] \leftarrow \text{Mem}(R[rs1] + \text{offset}, \text{word})$
addi rd, rs1, imm		0x13	0x0		$R[rd] \leftarrow R[rs1] + \text{imm}$
slli rd, rs1, imm			0x1	0x00	$R[rd] \leftarrow R[rs1] \ll \text{imm}$
slti rd, rs1, imm			0x2		$R[rd] \leftarrow (R[rs1] < \text{imm}) ? 1 : 0$
xori rd, rs1, imm			0x4		$R[rd] \leftarrow R[rs1] \wedge \text{imm}$
srli rd, rs1, imm			0x5	0x00	$R[rd] \leftarrow R[rs1] \gg \text{imm}$
srai rd, rs1, imm			0x5	0x20	$R[rd] \leftarrow R[rs1] \gg \text{imm}$
ori rd, rs1, imm			0x6		$R[rd] \leftarrow R[rs1] \vee \text{imm}$

INSTRUCTION	TYPE	OPCODE	FUNCT3	FUNCT7/IMM	OPERATION
<code>andi rd, rs1, imm</code>			0x7		$R[rd] \leftarrow R[rs1] \& imm$
<code>sw rs2, offset(rs1)</code>	S	0x23	0x2		$Mem(R[rs1] + offset) \leftarrow R[rs2]$
<code>beq rs1, rs2, offset</code>	SB	0x63	0x0		if( $R[rs1] == R[rs2]$ ) $PC \leftarrow PC + \{offset, 1b'0\}$
<code>blt rs1, rs2, offset</code>			0x4		if( $R[rs1]$ less than $R[rs2]$ (signed)) $PC \leftarrow PC + \{offset, 1b'0\}$
<code>bltu rs1, rs2, offset</code>			0x6		if( $R[rs1]$ less than $R[rs2]$ (unsigned)) $PC \leftarrow PC + \{offset, 1b'0\}$
<code>bne rs1, rs2, offset</code>			0x1		if( $R[rs1] != R[rs2]$ ) $PC \leftarrow PC + \{offset, 1b'0\}$
<code>lui rd, offset</code>	U	0x37			$R[rd] \leftarrow \{offset, 12b'0\}$
<code>jal rd, imm</code>	UJ	0x6f			$R[rd] \leftarrow PC + 4$ $PC \leftarrow PC + \{imm, 1b'0\}$
<code>jalr rd, rs, imm</code>	I	0x67	0x0		$R[rd] \leftarrow PC + 4$ $PC \leftarrow R[rs] + \{imm\}$

### 3) Controls

You can probably guess that control signals will play a very large part in this project. Figuring out all of the control signals may seem intimidating. We suggest taking a look at Discussion 6 to get started, and to remember that there is not a definitive set of control signals—walk through the datapath with different types of instructions, and when you see a mux or other component think about what selector/enable value you will need for that instruction.

Additionally, implementing the control signals can be done in many ways, including implementing the corresponding truth tables and using comparators. Since you are welcome to use any built-in Logisim circuits, we suggest using whichever component makes the most sense to you, whether performing logical operations on instruction bits and/or comparing fields to certain values.

Before you go and add the pipeline stages, we suggest you to verify the correctness of the control signals with the single-cycle CPU tests we give you.

### 4) Pipelining

Your processor will have a 2-stage pipeline:

1. **Instruction Fetch:** An instruction is fetched from the instruction memory. (Note: while you can, please do not calculate jump address in this stage. Instead, you should try to deal with the jump control hazard.)
2. **Execute:** The instruction is decoded, executed, and committed (written back). This is a combination of the remaining stages of a normal five-stage RISC-V pipeline.

First, make sure you understand what hazards you will have to deal with.

**Our ISA does not expose branch delay slots to software.** This means that the instruction immediately after a branch or jump is not executed if the branch is taken. This makes your task a bit more complex. By the time you have figured out that a branch or jump is in the execute stage, you have already accessed the instruction memory and pulled out (possibly) the wrong instruction. You will therefore need to "kill" instructions that are being fetched if the instruction under execution is a jump or a taken branch. **Instruction kills for this project MUST be accomplished by MUXing a `nop` into the instruction stream and sending the `nop` into the Execute stage instead of using the fetched instruction.** Notice that `0x00000000` is a `nop` instruction; please use this, as it will simplify grading and testing. You should only kill if a branch is taken (do not kill otherwise), but do kill on every type of jump.

Do not solve this issue by calculating branch offsets in the IF stage. Because we test your output against the reference every cycle, and the reference returns a `nop`, while it may be a conceptually correct solution, this will cause you to fail our tests.

Because all of the control and execution is handled in the Execute stage, **your processor should be more or less indistinguishable from a single-cycle implementation, barring the one-cycle startup latency and the branch/jump delays.** However, we will be enforcing the two-pipeline design. If you are unsure about pipelining, it is perfectly fine (maybe even recommended) to first implement a single-cycle processor. This will allow you to first verify that your instruction decoding, control signals, arithmetic

operations, and memory accesses are all working properly. From a single-cycle processor you can then split off the Instruction Fetch stage with a few additions and a few logical tweaks. Some things to consider:

- Will the IF and EX stages have the same or different `pc` values?
- Do you need to store the `pc` between the pipelining stages?
- To MUX a `nop` into the instruction stream, do you place it *before* or *after* the instruction register?
- What address should be requested next while the EX stage executes a `nop`? Is this different than normal?

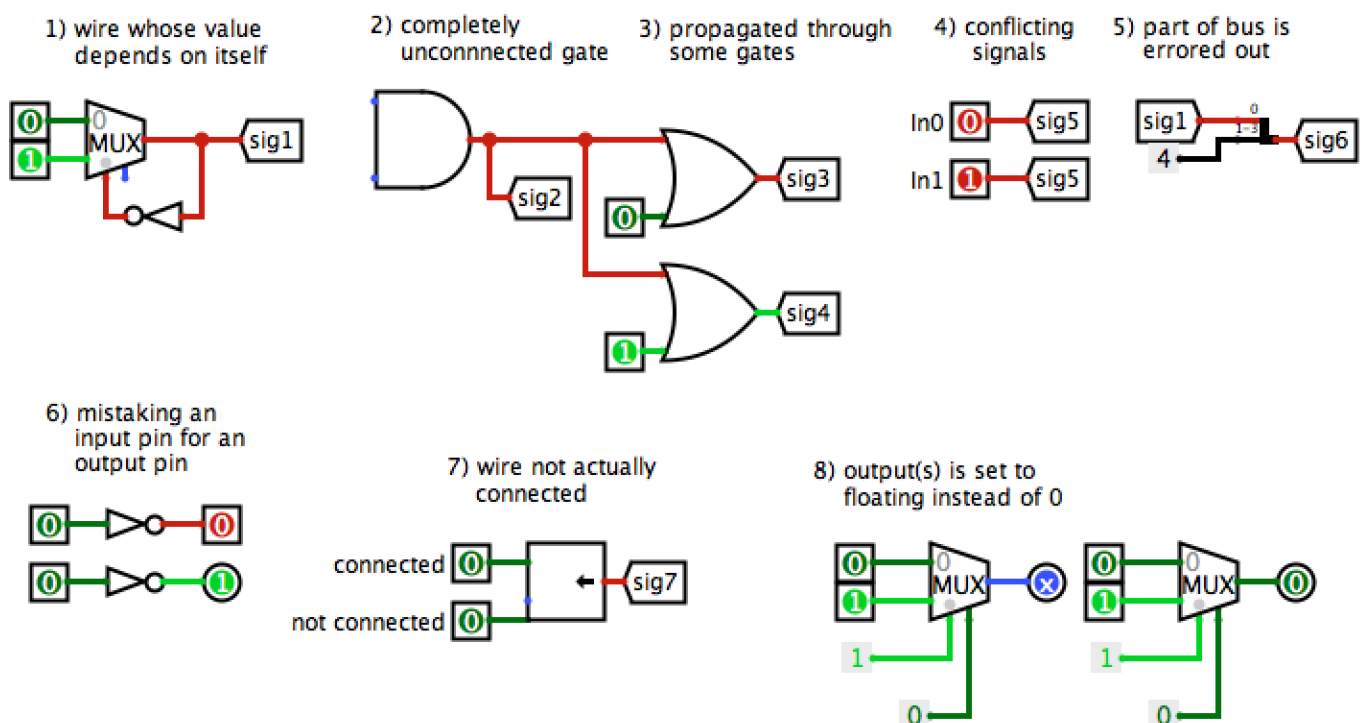
You might also notice a bootstrapping problem here: during the first cycle, the instruction register sitting between the pipeline stages won't contain an instruction loaded from memory. How do we deal with this? It happens that Logisim automatically sets registers to zero on reset; the instruction register will then contain a `nop`. We will allow you to depend on this behavior of Logisim. Remember to go to Simulate --> Reset Simulation (Ctrl+R) to reset your processor.

## Logisim Notes

If you are having trouble with Logisim, **RESTART IT and RELOAD your circuit!** Don't waste your time chasing a bug that is not your fault. However, if restarting doesn't solve the problem, it is more likely that the bug is a flaw in your project. Please post to Piazza about any crazy bugs that you find and we will investigate.

### Things to Look Out For

- Do **NOT** gate the clock! This is very bad design practice when making real circuits, so we will discourage you from doing this by heavily penalizing your project if you gate your clock.
- **BE CAREFUL with copying and pasting from different Logisim windows.** Logisim has been known to have trouble with this in the past.
- When you import another file (Project --> Load Library --> Logisim Library...), it will appear as a folder in the left-hand viewing pane. The skeleton files should have already imported necessary files.
- Changing attributes *before* placing a component changes the default settings for that component. So if you are about to place many 16-bit pins, this might be desirable. If you only want to change that particular component, place it first before changing the attributes.
- When you change the inputs & outputs of a sub-circuit that you have already placed in `main`, Logisim will automatically add/remove the ports when you return to `main` and this sometimes shifts the block itself. If there were wires attached, Logisim will do its automatic moving of these as well, which can be extremely dumb in some cases. Before you change the inputs and outputs of a block, it can sometimes be easier to first disconnect all wires from it.
- Error signals (red wires) are obviously bad, but they tend to appear in complicated wiring jobs such as the one you will be implementing here. It's good to be aware of the common causes while debugging:



Lastly, before you go into implementation, one last piece of advice: Modularize, modularize, modularize!

## Logisim's Combinational Analysis Feature

Logisim offers some functionality for automating circuit implementation given a truth table, or vice versa. Though not disallowed (enforcing such a requirement is impractical), use of this feature is discouraged. Remember that you will not be allowed to have a laptop running Logisim on the final.

---

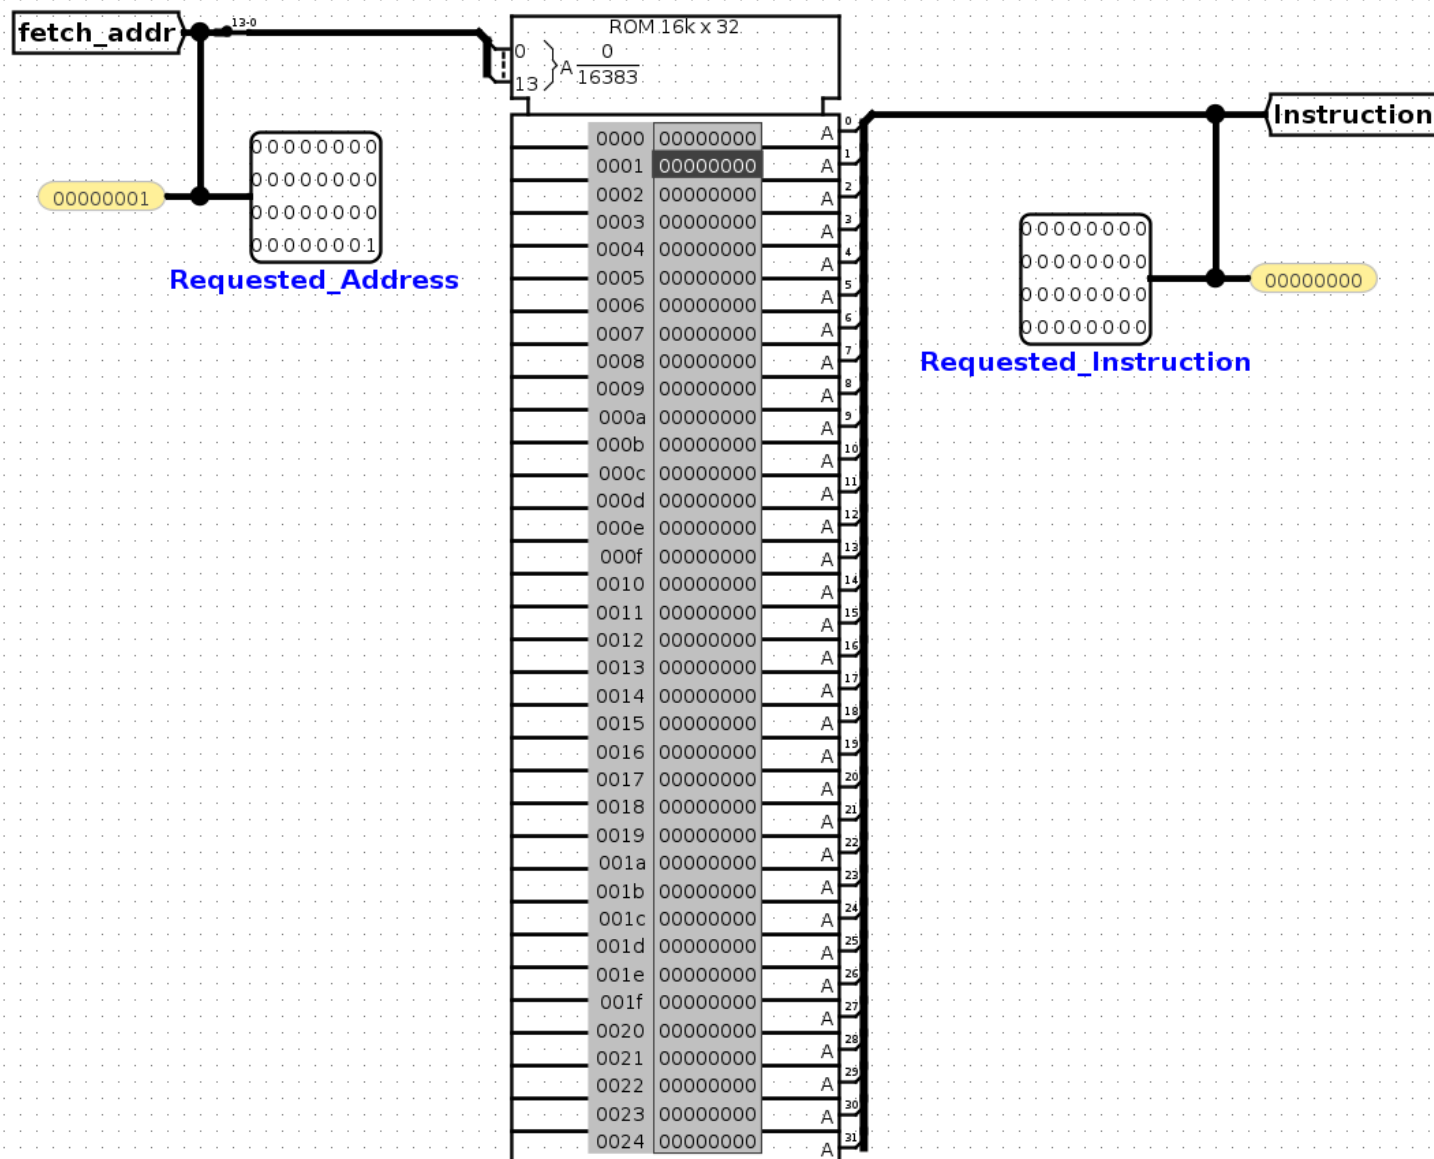
## Testing

For part 2, it is somewhat difficult to provide small unit tests such as the ones from part 1 since you are completing the full datapath. We recommend that you write short RISC-V programs to exercise your datapath, and loading them into a copy of *run.circ*, enabling simulation, and ticking through instruction by instruction to ensure that they work properly. You can use [Venus](#) to convert RISC-V program into Machine code (Under the *Simulation* tab, click on the *Dump* button to copy the machine code to clipboard). You can do this for both single-cycle and pipelined testing. You can also use the Trace Generator option to produce the expected output, by specifying which registers you want included in the trace options. Remember that the output is the 8 registers we implement besides x0, in the order seen in *run.circ*, the requested address and the returned instruction, and finally the test number. All of the values besides the test number are 32 bits/8 hex digits, and the test number is 16 bits/4 hex digits. However, because of the way branches work, you will only be able to use the trace for single-cycle tests. (Note: The output file in this part of the project is called "output", not "student\_output"). Please ask questions on Piazza if this is unclear--testing is very important, as the sanity tests are NOT comprehensive.


Once you have generated the machine code, you'll have to load it into the instruction memory unit *run.circ* and begin execution. To do so, first open *run.circ* and locate the Instruction Memory Unit.



## Instruction Memory



Click on the memory module and then, in the left sidebar, click on the "(Click to edit)" option next to "Contents". This will bring up a hex editor with the option to open a previously created hex file. This is where we load the file outputted by the assembler earlier.



Once you've loaded the machine code you can tick the clock manually and watch your CPU execute your program! You can double click on the CPU using the poke tool to take a look at how your datapath is behaving under the given input. You can compare the behavior of your CPU to the output of Venus emulator.

## Sanity tests

### a) Single-Cycle CPU

It is good practice to build a working single-cycle CPU before adding pipeline stages to it. The test harness for the single-cycle CPU is given to you in *tests/cpu\_single*. You can run the single cycle sanity test using the following command from your project folder:

```
./cpu-sanity-test-single.sh
```

This will copy your *cpu.circ*, *mem.circ*, *alu.circ* and *regfile.circ* into the *tests/cpu\_single* folder and then run the tests specified in *sanity\_test.py*.

### b) Two-Cycle CPU

You can run the pipelined sanity test using the following command from your project folder:

```
./cpu-sanity-test-pipelined.sh
```

The output of your cpu will be dumped in the *tests/cpu\_pipelined/output* folder. You can examine the output with the *binary\_to\_hex.py* script by:

```
cd tests/cpu_pipelined
python binary_to_hex.py output/CPU-<test_name>.out
```

To see what happened during a test, you can go to *tests/cpu\_pipelined* and open the *CPU-<test\_name>.circ* files. These are harnesses that are modified to have testing inputs in them. You can tick the clock and see the outputs of your circuit and compare them to what you would expect.



### c) Logisim Command-Line

If you are interested in how the circuits are being run through command-line, read the [Logisim Command-line Verification](#) page. Running the following command under *tests/cpu-pipelined/* will load the circuit *CPU-<test\_name>.circ* and execute the test in logisim:

```
java -jar logisim.jar CPU-<test_name>.circ -tty table
```

Logisim loads the current values of the output pins; if any have changed from the previous propagation, then all values are displayed in tab-delimited format.

## Submission

### Files to submit

```
part2/cpu.circ
part2/alu.circ
part2/regfile.circ
```

We will be using our own versions of the *run.circ* files with the test cases loaded into the memory, so you do not need to submit those. In addition, you should NOT depend on any changes you make to those files.

### Submission Instructions

There are **two** steps required to submit proj3. Failure to perform both steps will result in loss of credit:

1. First, you must submit glookup on the instructional servers. This assumes that you followed the earlier instructions and did all of your work inside of your git repository. To submit, follow these instructions after logging into your -XXX class account:

```
cd ~/proj3                                # Or wherever your proj3 git repo is
submit proj3
```

Once you type `submit proj3`, follow the prompts generated by the submission system. It will tell you when your submission has been successful and you can confirm this by looking at the output of `glookup -t`. You should be submitting 3 files: *alu.circ*, *regfile.circ*, and *cpu.circ*. We will run the autograder tests twice, once with our staff solution *alu.circ* and *regfile.circ*, and once with your *alu.circ* and *regfile.circ*, and your score will be the maximum of the two runs.

2. Additionally, you must submit proj3 to your Bitbucket repository:

```
cd ~/proj3                                # Or wherever your proj3 git repo is
git add -u
git commit -m "project 3 submission"        # The commit message doesn't have to match exactly.
git tag "proj3-sub"                        # The tag MUST be "proj3-sub". Failure to do so will result in loss of credit.
git push origin proj3-sub                  # This tells git to push the commit tagged proj3-sub
```

### Resubmitting

If you need to re-submit, you can follow the same set of steps that you would if you were submitting for the first time, but you will need to use the `-f` flag to tag and push to Bitbucket:

```
# Do everything as above until you get to tagging
git tag -f "proj3-sub"
git push -f origin proj3-sub
```

Note that in general, force pushes should be used with caution. They will overwrite your remote repository with information from your local copy. As long as you have not damaged your local copy in any way, this will be fine.

## Grading

This project will be graded in large part by an automatic grading script. For regrades, we will try to look to see if there is a simple wiring problem. If we can find one, we will give you the new score from the autograder minus a deduction based on the severity of the wiring problem. For this reason, neatness will play a factor in wiring – please try to make your circuits neat and readable. (Neat circuits will also make it much easier for the staff to help you debug!)